

Simulating Acoustics with Ray-Tracing in Rust

Kristofer Rye

January 27, 2020

Abstract

Ray tracing, a common method used in computer graphics to produce photo-realistic images, has natural applications to the study of acoustics. Sound moving through space can be modeled by particles which represent the wavefront as it moves through space, and sufficient numbers of particles can produce a well-resolved profile of received sound that is emitted from that point. We implement a primitive simulation system to approximate the acoustics of a modeled space and examine its resulting characteristics.

1 Introduction

Ray tracing is a common method used in the field of computer graphics to simulate the motion of light through space, and is commonly used to produce photo-realistic images. The idea of ray tracing in computer graphics originally came from the problem of computing the projection of shadows, [1] but developed in both accuracy and speed after it was popularized through its use in computer graphics. Originally conceived in the field of optics [4], ray tracing eventually reached popularity in cases where it was practical. Ray tracing has also seen use in physics research settings, especially for modeling the propagation of seismic waves, radio signals, and underwater sound.

A few things about sound make ray tracing particularly useful in the modeling of acoustics. First, since sound in air moves at a velocity far less than that of light, this means its motion through the space can be accurately approximated with particles moving at non-relativistic speeds.

Rust is a programming language of rising prominence with a stated goal of “empowering everyone to build reliable and efficient software.” [3] Rust was selected for this project because of a few key distinguishing factors. These include Rust’s exceptional memory safety guarantees, the inherent fearlessness of concurrency of Rust code, and the fact that Rust can be compiled both to native machine code (on x86_64 and similar devices) as well as WebAssembly, which allows it to be embedded within the browser rather cleanly.

In addition, Rust’s type system, specifically its rich support for generic functions and generic types, means that all of the algorithms with which we are concerned can be implemented for all types that support the underlying operation

generics. For a specific example, this means that the same implementation of our intersection algorithm can work in both two and three dimensions, as it makes use of the dot product which is well-defined in both dimensions.

2 Algorithms

The system we developed allows for the measurement of the acoustics of a modeled space with a specifiable level of detail. Wavefronts are approximated as particles carrying frequency and amplitude data, and are emitted omni-directionally (uniformly distributed, and randomly) within the space, and objects are modeled as either spheres (with a given origin and radius) or collections of triangles.

Our simulation consists of a loop (Algorithm 1) that continues endlessly until all sounds emitted at the start of the simulation have either been received or would be imperceptibly weak by the receiver. For each iteration through the simulation loop, every sound ray is checked against every object in the scene for an intersection (“hit”), and the intersection results are stored in a container which automatically sorts the stored sounds in ascending order by time, such as a binary tree. As a result, the time complexity of checking every sound with every object has time complexity $\mathcal{O}(m \cdot n)$ where m denotes the number of sounds and n denotes the number of objects. The claim that this worst-case time complexity is the best possible without advanced techniques needs verification, but hopefully intuitively makes sense: in short, culling would only remove a certain proportion of the objects in the scene, but all objects still need to be checked against all sounds in order to find the earliest hit, since that is what would happen first.

The algorithms for identifying intersections of sound (approximated by rays) and surfaces (approximated by triangles and spheres) are taken from relevant literature sources. Notably, we use the fast and lightweight algorithm presented by Möller and Trumbore in [2], which can be used not only to determine *whether* a given ray intersects with a triangle, but also *exactly where*, and *at exactly what time*, both of which are relevant to our simulation. The ray-sphere intersection was derived using Wolfram Mathematica to generate a parameterized form of the sphere equation with one parameter (time), and this was then translated into code. Since a ray can intersect with a sphere at exactly zero, one, or two points, the intersection returns both points and the earlier time is selected. Unit

Algorithm 1 The simulation structure

```

Objects  $\leftarrow$  [room geometry]
Objects  $\leftarrow$  [receivers]
for emitter  $\in$  Emitters do
  Sounds  $\leftarrow$  emitter.emit
end for
repeat
  for sound  $\in$  Sounds do
    I  $\leftarrow$  []
    for object  $\in$  Objects do
      I  $\leftarrow$  sound.hit?(object)
    end for
    hit  $\leftarrow$  I.first
    if hit.object is a Receiver or
      sound.amplitude * hit.reflectance <  $\epsilon$ 
    then
      hit.object.hits  $\leftarrow$  hit
      Sounds.delete(hit.sound)
    else
      hit.sound.bounce(hit.object)
    end if
  end for
until Sounds =  $\emptyset$ 

```

normal vectors are generated in both intersection algorithms and are used by the scene simulator to compute in what direction the outgoing ray moves.

One limitation of our system is that it assumes perfectly elastic collisions with un-movable surroundings, and—perhaps more significantly—that sound exclusively bounces at a lower amplitude when it interacts with a surrounding. In reality, a wavefront interacts with the actual materials of the wall on a highly detailed basis, and the level of detail required proves to be impractical. Instead, our system requires that entire surfaces (groups of triangles) are approximated with a single coefficient representing how much of the sound they reflect, and as additional detail is needed, additional surfaces must be added. Furthermore, we assume exactly one medium with one speed of sound, which is not wholly realistic. A given space could have varying speeds of sound in different areas as factors like lighting induce thermal differences and air currents. The difference in sound profile is assumed to be negligible, but in practice might be somewhat significant.

An alternative or more primitive approach to implementation would have rays stepping forward an arbitrarily small amount, checking for bounces off each object at each time step, and moving in accordance with a *computed* speed of sound, rather than a constant speed of sound. Such an approach would allow for the simulator to account for differences in localized temperature and speed of sound, but since in most time steps no collisions would be observed, a large amount of time would be wasted checking for those collisions.

3 Results

The simulated space is a 20 meter long \times 10 meter wide \times 10 meter tall space, with a front wall, ceiling, and floor each absorbing 20% of the intensity of the sound that they receive, a back wall absorbing 80% of the sound that hits it, and left and right walls each absorbing 60% of the sound that hits them. The emitter is located 1 m towards the front wall, and the receiver is a 0.1 m-radius sphere, of comparable size to the human head. According to the textbook’s approximation for reverberation time, the reverberation time (RT) of this space would be given by

$$\begin{aligned} RT &= 0.161 \frac{V}{A} \\ &= 0.161 \frac{2000 \text{ m}^3}{(100 + 2 \cdot 200) \cdot .20 + 100 \cdot .80 + (2 \cdot 200) \cdot .60 \text{ m}^2} \\ &\approx .767 \text{ s} \end{aligned}$$

When the simulation is run, a configurable number of rays—which, for the sake of the included figures, is 1,000,000—is produced by the emitter, distributed uniformly along a unit sphere and then scaled by the speed of sound. The simulation is then run to completion, taking approximately 8 iterations before all sounds have been absorbed, which represents approximately 8 bounces

off of walls before the sound is received. The simulation subroutine then yields back a table containing the observed “hit” data (including what intensity was recorded) and time of the observation. These data can then be plotted using a tool like gnuplot with time on the horizontal axis and intensity on a logarithmic scale on the vertical axis, producing the following figures.

The first figure, Figure 1, shows one run of the simulation. Nearly all of the sound is received within the first two seconds, but the logarithmic scale demonstrates how since nearly all of the sound is reflected around only a few times, only three orders of magnitude on intensity are lost, representing a loss of only 30 dB of loudness. Likewise, the intensity levels observed are quantized onto a handful of levels that can be seen as horizontal groupings. Since the geometry of the space is incredibly simple, each intensity level corresponds to a certain, countable and finite number of bounces off of a surface, and there is not much variability as there are only six surfaces with which a sound wave can interact with.

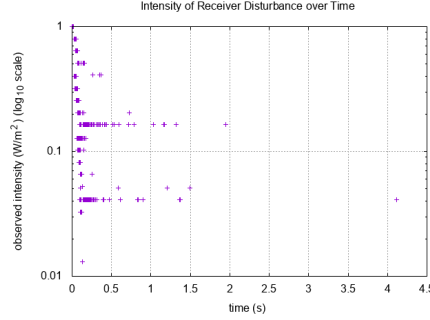


Figure 1: Simulation results




Figure 2: Simulation results, constrained to $[0, 2)$ seconds

a bounce off the back wall, with varying angles of incidence corresponding to a longer time delay. Adjusting the reflectance parameters of the ceiling and floor produces other groupings like these at different intensity levels, confirming how they are formed. Indeed, most of the sound is observed to dissipate by 0.767

Figure 2 displays the same results, constrained to only the first two seconds on the horizontal axis. Here various groupings of sound waves can be observed, the first of highest intensity from a time near zero seconds to a time near 0.1 seconds, and with the second highest intensities coming shortly after that. There is also a significant group of particles which reached the receiver with an intensity between 0.1 and 0.2 W/m^2 , corresponding to a bounce off the ceiling and the floor, or the front wall and either the floor or the ceiling. An additional long group of particles which reached the receiver with roughly 0.03 W/m^2 corresponding to

seconds, in accordance with the prediction by the textbook.

3.1 Limitations

Our system does not support filtering different frequencies at different levels, however it would be trivial to extend it to work as such. When computing the resulting amplitude, instead of simply multiplying by a constant, a more complicated calculation could be performed based on linear interpolation over known reflectances at known frequencies, for example. We also assume the same amount is absorbed by a surface with a blow by a wave compared to a head-on wavefront, which is not necessarily the case in optical ray tracing. Furthermore, as mentioned before, our simulation only runs once with one wavefront, simulating the sound of a sharp change in pressure, such as a piercing explosion, at 120 dB. This does not model most musical sounds, which would consist of a fundamental and its partial series based on the timbre of the instrument. Further modeling would be required to describe how the timbre of transmitted sound changes as it goes through the space.

4 Conclusion

We have implemented a ray-tracing acoustics engine that produces somewhat realistic results in a simple but still complete example. Our system allows for maximum configurability and uses relatively fast technologies so that its level of detail can still remain high as more complex simulations are performed. Indeed, ray tracing proves very promising for validating assumptions about spatial acoustics in the design phase rather than in the implementation phase of construction, which means that ray tracing has a lot of potential uses in the study of acoustics as it relates to architectural design.

References

- [1] Arthur Appel. “Some techniques for shading machine renderings of solids”. In: *Proceedings of the April 30–May 2, 1968, spring joint computer conference on - AFIPS ‘68 (Spring)*. ACM Press, 1968. DOI: 10.1145/1468075.1468082. URL: <https://doi.org/10.1145/1468075.1468082>.
- [2] Tomas Möller and Ben Trumbore. “Fast, Minimum Storage Ray–Triangle Intersection”. In: *Journal of Graphics Tools* 2.1 (Jan. 1997), pp. 21–28. DOI: 10.1080/10867651.1997.10487468. URL: <https://doi.org/10.1080/10867651.1997.10487468>.
- [3] *Rust Programming Language*. URL: <https://www.rust-lang.org>.
- [4] G. H. Spencer and M. V. R. K. Murty. “General Ray-Tracing Procedure†”. In: *J. Opt. Soc. Am.* 52.6 (June 1962), pp. 672–678. DOI: 10.1364/JOSA.52.000672. URL: <http://www.osapublishing.org/abstract.cfm?URI=josa-52-6-672>.